

HW1

I. Compiling the Linux Kernel

參考: <https://phoenixnap.com/kb/build-linux-kernel>

1. 根據助教提供之設定步驟建立環境。

2. building the kernel 前，安裝必要的 package

Command: **sudo apt-get install git build-essential ncurses-dev libssl-dev bc flex libelf-dev bison**

3. 修改 .config, 然後設定作業中所要求的 kernel local version suffix

Command: **make menuconfig**

vim .config

```
CONFIG_CC_CAN_LINK=y
CONFIG_CC_CAN_LINK_STATIC=y
CONFIG_CC_HAS_ASM_GOTO_OUTPUT=y
CONFIG_CC_HAS_ASM_GOTO_TIED_OUTPUT=y
CONFIG_CC_HAS_ASM_INLINE=y
CONFIG_CC_HAS_NO_PROFILE_FN_ATTR=y
CONFIG_PAHOLE_VERSION=0
CONFIG_IRQ_WORK=y
CONFIG_BUILDTIME_TABLE_SORT=y
CONFIG_THREAD_INFO_IN_TASK=y

#
# General setup
#
CONFIG_INIT_ENV_ARG_LIMIT=32
# CONFIG_COMPILE_TEST is not set
# CONFIG_WERROR is not set
CONFIG_LOCALVERSION="-os-312512032"
# CONFIG_LOCALVERSION_AUTO is not set
CONFIG_BUILD_SALT=""
CONFIG_HAVE_KERNEL_GZIP=y
CONFIG_HAVE_KERNEL_BZIP2=y
CONFIG_HAVE_KERNEL_LZMA=y
".config" 11833L, 272557B                               31,34          0%
```

4. building the kernel 、安裝 modules

Command: **make**

sudo make modules_install

5. 確認要安裝的 kernel 版本

Command: **make kernelrelease**

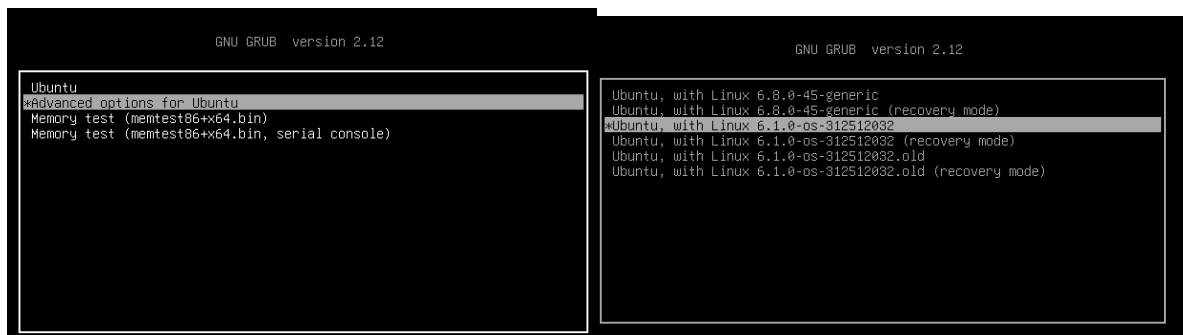
```
benson-312512032@benson-312512032-VirtualBox:~/linux$ make kernelrelease
6.1.0-os-312512032
```

6. 安裝 kernel.

Command: **sudo make install**

7. 設定開機時所要選的 kernel.(開機時按 Esc 進入 GNU GRUB 確認你的 kernel 在選單的哪裡

Example: Advanced options for Ubuntu > Ubuntu, with Linux 6.1.0-os-312512032, 更新 /etc/default/grub 的 GRUB_DEFAULT)



```
# If you change this file, run 'update-grub' afterwards to update
# /boot/grub/grub.cfg.
# For full documentation of the options in this file, see:
# info -f grub -n 'Simple configuration'

GRUB_DEFAULT="1>2"
GRUB_TIMEOUT_STYLE=hidden
GRUB_TIMEOUT=0
GRUB_DISTRIBUTOR="( . /etc/os-release; echo ${NAME:-Ubuntu} ) 2>/dev/null || ech
o Ubuntu"
GRUB_CMDLINE_LINUX_DEFAULT="quiet splash"
GRUB_CMDLINE_LINUX=""

# If your computer has multiple operating systems installed, then you
# probably want to run os-prober. However, if your computer is a host
# for guest OSes installed via LVM or raw disk devices, running
# os-prober can cause damage to those guest OSes as it mounts
# filesystems to look for things.
GRUB_DISABLE_OS_PROBER=false

# Uncomment to enable BadRAM filtering, modify to suit your needs
# This works with Linux (no patch required) and with any kernel that obtains
# the memory map information from GRUB (GNU Mach, kernel of FreeBSD ...)
"/etc/default/grub" 40L, 1553B 1,1 Top
```

8. 更新 GRUB bootloader

Command: **sudo update-grub**

```
benson-312512032@benson-312512032-VirtualBox:~/linux$ sudo update-grub
Sourcing file '/etc/default/grub'
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-6.8.0-45-generic
Found initrd image: /boot/initrd.img-6.8.0-45-generic
Found linux image: /boot/vmlinuz-6.1.0-os-312512032
Found initrd image: /boot/initrd.img-6.1.0-os-312512032
Found linux image: /boot/vmlinuz-6.1.0-os-312512032.old
Found initrd image: /boot/initrd.img-6.1.0-os-312512032
Found memtest86+x64 image: /boot/memtest86+x64.bin
Warning: os-prober will be executed to detect other bootable partitions.
Its output will be used to detect bootable binaries on them and create new boot entries.
Adding boot menu entry for UEFI Firmware Settings ...
done
```

9. 重新開機

Command: **sudo reboot**

10. 確認當前之 kernel 版本

Command: **uname -a**

cat /etc/os-release

```
benson-312512032@benson-312512032-VirtualBox:~/Desktop$ uname -a
Linux benson-312512032-VirtualBox 6.1.0-os-312512032 #2 SMP PREEMPT_DYNA
MIC Tue Sep 24 15:37:40 CST 2024 x86_64 x86_64 x86_64 GNU/Linux
benson-312512032@benson-312512032-VirtualBox:~/Desktop$ cat /etc/os-rele
ase
PRETTY_NAME="Ubuntu 24.04.1 LTS"
NAME="Ubuntu"
VERSION_ID="24.04"
VERSION="24.04.1 LTS (Noble Numbat)"
VERSION_CODENAME=noble
ID=ubuntu
ID_LIKE=debian
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/priv
acy-policy"
UBUNTU_CODENAME=noble
LOGO=ubuntu-logo
benson-312512032@benson-312512032-VirtualBox:~/Desktop$
```

II. Implementing a new System Calls

1. 在 linux 內建立 revstr 資料夾,並撰寫該 system call 的內容於 revstr.c

此函數有兩個參數 msg、len，因此使用 SYSCALL_DEFINE2.

由於 user/kernel space 的資料無法直接與 kernel/user space 共用,因此使用 copy_from_user()、copy_to_user,進行資料的傳遞,並以動態記憶體配置的方式給儲存資料(為什麼是 len+1,因為 user space 傳遞過來的長度不包含 '\0',導致%s 在印 System 時會多一串亂碼,因此透過此方式解決),接著由頭尾互相交換直到中間,完成字串反轉,最後釋放動態記憶體。

```
#include <linux/kernel.h>
#include <linux/syscalls.h>
#include <linux/slab.h>
#include <linux/uaccess.h>
// syscall_num =451
SYSCALL_DEFINE2(revstr, char __user * ,msg,unsigned int ,len){
    //printk(KERN_INFO"%d",len);
    char* kernel_msg = kmalloc((len+1)*sizeof(char),GFP_KERNEL);
    if(! kernel_msg){
        printk(KERN_ERR "memory allocation failed \n");
        return -ENOMEM ;
    }
    if (copy_from_user(kernel_msg,msg,len)){
        printk(KERN_ERR "can not copy form user ,please check your address!\n");
        kfree(kernel_msg);
        return -EFAULT;
    }
    /*****revstr function*****/
    kernel_msg[len]='\0';
    printk(KERN_INFO "The origin string: %s\n",kernel_msg);
    for(unsigned int i=0,j=len-1 ;i<j;i++,j--){
        char temp = kernel_msg[i];
        kernel_msg[i]=kernel_msg[j];
        kernel_msg[j]=temp;
    }
    kernel_msg[len]='\0';
    printk(KERN_INFO "The reversed string: %s\n",kernel_msg);
    /*****
    if ( copy_to_user(msg,kernel_msg,len)){
        printk(KERN_ERR "can not copy to user!!\n");
        return -EFAULT;
    }
    kfree(kernel_msg);
    return 0;
    *****/
}
```

2. 撰寫編譯 revstr.c 的 Makefile 於 ~/linux/revstr,使其被編譯並連結至 kernel image 內

```
obj-y := revstr.o
```

3. 將所撰寫的 sys 加入 ~/linux/arch/x86/entry/syscalls/syscall_64.tbl，加入 64 bit 的最後面，我們的 system call number 是 451，名稱為 revstr,進入點為 sys_revstr

```
# don't use numbers 387 through 423, add new calls after the last
# 'common' entry
424 common pidfd_send_signal sys_pidfd_send_signal
425 common io_uring_setup sys_io_uring_setup
426 common io_uring_enter sys_io_uring_enter
427 common io_uring_register sys_io_uring_register
428 common open_tree sys_open_tree
429 common move_mount sys_move_mount
430 common fsopen sys_fsopen
431 common fsconfig sys_fsconfig
432 common fsmount sys_fsmount
433 common fspick sys_fspick
434 common pidfd_open sys_pidfd_open
435 common clone3 sys_clone3
436 common close_range sys_close_range
437 common openat2 sys_openat2
438 common pidfd_getfd sys_pidfd_getfd
439 common faccessat2 sys_faccessat2
440 common process_madvise sys_process_madvise
441 common epoll_pwait2 sys_epoll_pwait2
442 common mount_setattr sys_mount_setattr
443 common quotactl_fd sys_quotactl_fd
444 common landlock_create_ruleset sys_landlock_create_ruleset
445 common landlock_add_rule sys_landlock_add_rule
446 common landlock_restrict_self sys_landlock_restrict_self
447 common memfd_secret sys_memfd_secret
448 common process_mrelease sys_process_mrelease
449 common futex_waitv sys_futex_waitv
450 common set_mempolicy_home_node sys_set_mempolicy_home_node
451 common revstr sys_revstr
#
# Due to a historical design error, certain syscalls are numbered differently
# in x32 as compared to native x86_64. These syscalls have numbers 512-547.
```

4. 於 `~/linux/include/linux/syscalls.h` 宣告 `sys_revstr` 的原型，使 Linux Kernel 編譯器就能識別這個新系統呼叫，並在需要的地方呼叫它

```
/* obsolete: ipc */
asmlinkage long sys_ipc(unsigned int call, int first, unsigned long
                        unsigned long third, void __user *ptr, long fifth);

/* obsolete: mm/ */
asmlinkage long sys_mmap_pgoff(unsigned long addr, unsigned long
                                unsigned long prot, unsigned long flags,
                                unsigned long fd, unsigned long pgoff);
asmlinkage long sys_old_mmap(struct mmap_arg_struct __user *arg);

/*
 * Not a real system call, but a placeholder for syscalls which are
 * not implemented -- see kernel/sys_ni.c
 */
asmlinkage long sys_ni_syscall(void);
/* revstr/revstr.h */
asmlinkage long sys_revstr(char __user * msg ,unsigned int len);
```

5. 將資料夾 `~/linux/revstr` 之相對路徑加入 `~/linux/Makefile` 中的 `core-y`,使其加入在編譯過程中

```
ifneq ($(filter all modules nsdeps %compile_commands.json clang-%,$(MAKECMDGOALS)),)
    KBUILD_MODULES := 1
endif

ifneq ($(MAKECMDGOALS),)
    KBUILD_MODULES := 1
endif

export KBUILD_MODULES KBUILD_BUILTIN

ifdef need-config
include include/config/auto.conf
endif

ifneq ($(KBUILD_EXTMOD),)
# Objects we will link into vmlinux / subdirs we need to visit
core-y      := revstr/
drivers-y    :=
libs-y      := lib/
endif # KBUILD_EXTMOD
```

6. 重新燒入並重新啟動

Command: `make -j8 && sudo make modules_install && sudo make install && sudo reboot`

7. 編寫測試檔案，`__NR_revstr` 是 system call number,藉此讓 `syscall()` function 執行 system call table 中對應號碼的 system call

```
#include <unistd.h>
#include <string.h>
#include <stdio.h>
#include <assert.h>

#define __NR_revstr 451

int main(int argc ,char* argv[]) {

    char str1[20] = "hello";
    printf("Ori: %s\n",str1);
    int ret1 = syscall(__NR_revstr,str1,strlen(str1));
    assert(ret1 == 0);
    printf("Rev: %s\n",str1);

    char str2[20] = "Operating System";
    printf("Ori: %s\n",str2);

    int ret2=syscall(__NR_revstr,str2,strlen(str2));
    assert(ret2 == 0);
    printf("Rev: %s\n",str2);

    return 0;
}
```

8. 執行測試檔案和查看 user space 的結果及 dmesg 之 kernel space 的結果

```
benson-312512032@benson-312512032-VirtualBox:~/Desktop/HW1$ ./test_
revstr
Ori: hello
Rev: olleh
Ori: Operating System
Rev: metsyS gnitarep0
```

```
[ 139.913654] The origin string: hello
[ 139.913657] The reversed string: olleh
[ 139.913661] The origin string: Operating System
[ 139.913662] The reversed string: metsyS gnitarep0
```